

# Lab8

Paulina Duda

```
[104]: import numpy as np
import tensorflow.keras as keras
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras import models
from tensorflow.keras.utils import to_categorical
from tensorflow.keras import backend as K
```

HOMEWORK 1 Build a classifier for fashion MNIST.

1. Use exactly the same architectures (both densely connected layers and from convolutional layers) as the above MNIST e.g., replace the dataset. Save the Jupyter Notebook in its original format and output a PDF file after training, testing, and validation. Make sure to write down how do they perform (training accuracy, testing accuracy).
2. Improve the architecture. Experiment with different numbers of layers, size of layers, number of filters, size of filters. You are required to make those adjustment to get the highest accuracy. Watch out for overfitting – we want the highest testing accuracy! Please provide a PDF file of the result, the best test accuracy and the architecture (different numbers of layers, size of layers, number of filters, size of filters)

```
[107]: (train_images, train_labels), (test_images, test_labels) = keras.datasets.
↳fashion_mnist.load_data()

model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
model.add(layers.Dense(10, activation='softmax'))
model.compile(optimizer='rmsprop',
loss='mean_squared_error',
metrics=['accuracy'])

train_images_flat = train_images.reshape((60000, 28 * 28))
train_images_flat = train_images_flat.astype('float32') / 255
test_images_flat = test_images.reshape((10000, 28 * 28))
test_images_flat = test_images_flat.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

```
train_images.reshape((60000,28*28)).shape
```

```
model.fit(train_images_flat, train_labels, epochs=10, batch_size=128)
```

Train on 60000 samples

Epoch 1/10

60000/60000 [=====] - 1s 21us/sample - loss: 0.0317 - accuracy: 0.7799

Epoch 2/10

60000/60000 [=====] - 1s 18us/sample - loss: 0.0207 - accuracy: 0.8582

Epoch 3/10

60000/60000 [=====] - 1s 19us/sample - loss: 0.0185 - accuracy: 0.8724

Epoch 4/10

60000/60000 [=====] - 1s 19us/sample - loss: 0.0170 - accuracy: 0.8834

Epoch 5/10

60000/60000 [=====] - 1s 19us/sample - loss: 0.0164 - accuracy: 0.8880

Epoch 6/10

60000/60000 [=====] - 1s 19us/sample - loss: 0.0155 - accuracy: 0.8961

Epoch 7/10

60000/60000 [=====] - 1s 21us/sample - loss: 0.0150 - accuracy: 0.8994

Epoch 8/10

60000/60000 [=====] - 1s 19us/sample - loss: 0.0145 - accuracy: 0.9017

Epoch 9/10

60000/60000 [=====] - 1s 20us/sample - loss: 0.0141 - accuracy: 0.9052

Epoch 10/10

60000/60000 [=====] - 1s 19us/sample - loss: 0.0136 - accuracy: 0.9089

[107]: <tensorflow.python.keras.callbacks.History at 0x7fec5413b040>


```
[109]: test_loss, test_acc = model.evaluate(test_images_flat, test_labels)
print(f'test accuracy: {test_acc}')
```

test accuracy: 0.8842999935150146

```
[110]: predictions = model.predict(train_images_flat[:13])
img_num = 4
print(predictions[img_num])
```

```
print(train_labels[img_num])
plt.matshow(train_images[img_num], cmap='gray')
plt.show()
```

```
[3.1008202e-01 2.7938515e-01 3.9449881e-04 2.6109755e-01 2.7240140e-04
 5.3294832e-07 1.4871316e-01 1.0060092e-07 5.4574037e-05 1.1384736e-09]
[1. 0. 0. 0. 0. 0. 0. 0. 0.]
```



```
[111]: model2 = models.Sequential()
model2.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.Flatten())
model2.add(layers.Dense(64, activation='relu'))
model2.add(layers.Dense(10, activation='softmax'))
train_images_conv = train_images.reshape((60000, 28, 28, 1))
train_images_conv = train_images_conv.astype('float32') / 255
test_images_conv = test_images.reshape((10000, 28, 28, 1))
test_images_conv = test_images_conv.astype('float32') / 255

model2.compile(optimizer='rmsprop',
loss='categorical_crossentropy',
```

```
metrics=['accuracy'])

model2.fit(train_images_conv, train_labels, epochs=8, batch_size=64)
```

Train on 60000 samples

Epoch 1/8

60000/60000 [=====] - 11s 190us/sample - loss: 0.5553 - accuracy: 0.7952

Epoch 2/8

60000/60000 [=====] - 11s 178us/sample - loss: 0.3363 - accuracy: 0.8776

Epoch 3/8

60000/60000 [=====] - 9s 153us/sample - loss: 0.2840 - accuracy: 0.8972

Epoch 4/8

60000/60000 [=====] - 10s 166us/sample - loss: 0.2521 - accuracy: 0.9083

Epoch 5/8

60000/60000 [=====] - 10s 160us/sample - loss: 0.2292 - accuracy: 0.9159

Epoch 6/8

60000/60000 [=====] - 10s 167us/sample - loss: 0.2105 - accuracy: 0.9227

Epoch 7/8

60000/60000 [=====] - 10s 171us/sample - loss: 0.1939 - accuracy: 0.9292

Epoch 8/8

60000/60000 [=====] - 10s 167us/sample - loss: 0.1807 - accuracy: 0.9336


[111]: <tensorflow.python.keras.callbacks.History at 0x7fec5466f250>

```
[114]: test_loss, test_acc = model2.evaluate(test_images_conv, test_labels)
       print(f'test accuracy: {test_acc}')
```

test accuracy: 0.9096999764442444

```
[115]: predictions = model2.predict(train_images_conv[:13])
       img_num = 8
       print(predictions[img_num])
       print(train_labels[img_num])
       plt.matshow(train_images[img_num], cmap='gray')
       plt.show()
```

```
[3.3107010e-09 7.5370151e-20 3.0884661e-14 5.8351337e-22 3.3076137e-17
 1.0000000e+00 3.4019350e-11 7.9130170e-13 5.5454773e-14 1.6382093e-13]
[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```



```
[116]: (train_images, train_labels), (test_images, test_labels) = keras.datasets.
        ↪fashion_mnist.load_data()

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

train_images.reshape((60000,28*28))

model2 = models.Sequential()
model2.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, ↪
        ↪1)))
model2.add(layers.Conv2D(128, (3, 3), activation='relu'))
model2.add(layers.MaxPooling2D((3, 3)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.Flatten())
model2.add(layers.Dense(64, activation='relu'))
model2.add(layers.Dense(10, activation='softmax'))

train_images_conv = train_images.reshape((60000, 28, 28, 1))
train_images_conv = train_images_conv.astype('float32') / 255
test_images_conv = test_images.reshape((10000, 28, 28, 1))
test_images_conv = test_images_conv.astype('float32') / 255
```

```
model2.compile(optimizer='rmsprop',
loss='categorical_crossentropy',
metrics=['accuracy'])

model2.fit(train_images_conv, train_labels, epochs=4, batch_size=64)
```

Train on 60000 samples

Epoch 1/4

60000/60000 [=====] - 38s 633us/sample - loss: 0.5127 - accuracy: 0.8138

Epoch 2/4

60000/60000 [=====] - 41s 677us/sample - loss: 0.3008 - accuracy: 0.8905

Epoch 3/4

60000/60000 [=====] - 40s 661us/sample - loss: 0.2525 - accuracy: 0.9083

Epoch 4/4

60000/60000 [=====] - 39s 658us/sample - loss: 0.2203 - accuracy: 0.9193

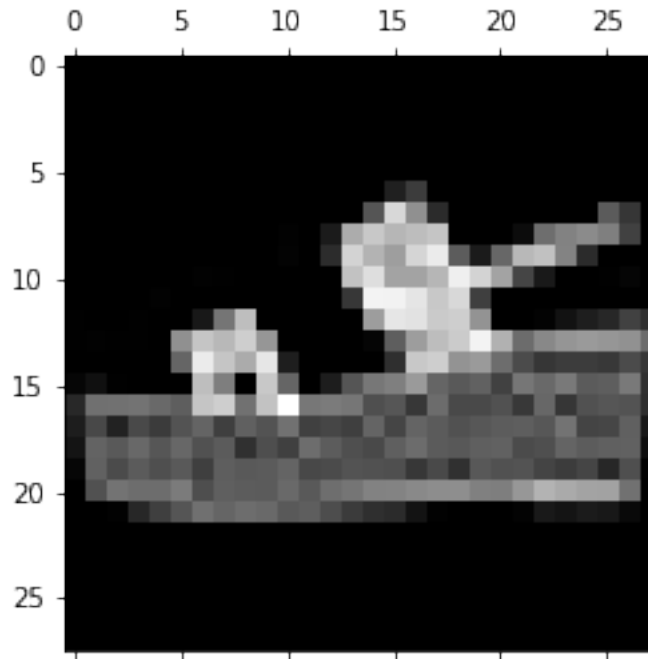
[116]: <tensorflow.python.keras.callbacks.History at 0x7fec9d3d2310>

```
[117]: test_loss, test_acc = model2.evaluate(test_images_conv, test_labels)
print(f'test accuracy: {test_acc}')
```

test accuracy: 0.9064000248908997

```
[118]: predictions = model2.predict(train_images_conv[:13])
img_num = 12
print(predictions[img_num])
print(train_labels[img_num])
plt.matshow(train_images[img_num], cmap='gray')
plt.show()
```

```
[2.5191978e-08 1.5257122e-10 1.6330008e-08 3.8754500e-08 1.3681105e-08
 9.9999249e-01 1.1286445e-08 4.7025646e-06 2.5639431e-06 1.0083376e-07]
[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```



## HOMEWORK 2:

Write a function that takes as arguments the name of the layer and filter index and outputs the displayable filter response. Then you can choose different filters and visualize which patterns they are responsive too! Submit the code (as Notebook) and at least 3 filter responses (a PDF file).

```
[119]: from tensorflow.keras.applications import VGG16
from tensorflow.keras import backend as K
import tensorflow as tf

tf.compat.v1.disable_eager_execution()
model = VGG16(weights='imagenet', include_top=False)

layer_name = 'block3_conv1'
filter_index = 0

layer_output = model.get_layer(layer_name).output
loss = K.mean(layer_output[:, :, :, filter_index])

grads = K.gradients(loss, model.input)[0]
iterate = K.function([model.input], [loss, grads])
loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])
```

```
[120]: model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_6 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		



```
[121]: input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.
step = 1.
for i in range(40):
    loss_value, grads_value = iterate([input_img_data])
    input_img_data += grads_value * step

#Postprocess to turn into displayable image


def deprocess_image(x):
    x -= x.mean()
    x /= (x.std() + 1e-5)
    x *= 0.1

    x += 0.5
    x = np.clip(x, 0, 1)

    x *= 255
    x = np.clip(x, 0, 255).astype('uint8')
    return x
```

```
[82]: plt.imshow(deprocess_image(input_img_data[0]))
```

```
[82]: <matplotlib.image.AxesImage at 0x7fec8853d6a0>
```



```
[95]: def displayable_response(layer_n, filter_i):

    layer_name = layer_n
    filter_index = filter_i

    layer_output = model.get_layer(layer_name).output
    loss = K.mean(layer_output[:, :, :, filter_index])

    grads = K.gradients(loss, model.input)[0]
    grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)

    iterate = K.function([model.input], [loss, grads])

    loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])

    input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.

    step = 1.

    for i in range(40):
        loss_value, grads_value = iterate([input_img_data])
        input_img_data += grads_value * step

    x = input_img_data[0]


    x -= x.mean()
    x /= (x.std() + 1e-5)
    x *= 0.1

    x += 0.5
    x = np.clip(x, 0, 1)


    x *= 255
    x = np.clip(x, 0, 255).astype('uint8')

    return x
```

```
[100]: plt.imshow(displayable_response('block2_conv2', 4))
plt.savefig("response1.pdf")
```



```
[101]: plt.imshow(displayable_response('block4_conv1', 8))
plt.savefig("response2.pdf")
```



```
[103]: plt.imshow(displayable_response('block4_pool', 10))  
plt.savefig("response3.pdf")
```

